

CONFIDENTIAL · Restricted to Quality and Release Stakeholders

Quality Tool-Agent Governance Summary Report

Quality Tool-Agent Governance Summary Report

Execution-type Agent Evaluation · Single-Indicator Statistics · L3 High Risk · Independent Execution

Fail

Fail

Tool-Agent Evaluation Gate **Fail**; Take one of five states from SanityOps Core Specification §6.2 Release Summary Status; Fail → FAIL **FAIL**。

Candidate: Reliability 50% below L3 threshold 100%, evaluation gate Fail, do not recommend release Voyager V3.

Report ID SO-QLT-TOOL-20260928	Evaluation Target Voyager	Evaluated Version V3
Risk Level L3 (High Risk · Irreversible/Sensitive Operations)	Execution Count 12 times	Success / Failed 6 / 6
Reliability 50%	Evaluation Gate Fail	Issuer / Time Quality Engineering Team · 09-28-2026

SanityOps Platform · Issued by Quality Engineering Team · Evaluation executed on 09-28-2026 · Report issued on 09-28-2026
This report is a summary; the full case list and per-failure execution evidence are in the same-batch Quality Tool-Agent Detailed Report (quality-tool-report).

Cover Glossary Quick Reference

Reliability Reliability = Success Count / Total Execution Count × 100% (Quality Tool-Agent White Paper §3.1.2 'Definition and Measurement of Reliability')	L1/L2/L3 Risk grading, thresholds 90% / 95% / 100% (Quality Tool-Agent White Paper §4.2.3 'Grading and Evaluation Rigor')	Gate $R \geq \text{threshold} \rightarrow \text{Pass}$, otherwise Fail (Quality Tool-Agent White Paper §4.3 'Evaluation Workflow' Step 6)
Ratchet Records the historical best; alerts when below it (Quality Tool-Agent White Paper §4.5.2 'Mechanism Principle')	Sandbox A production-equivalent, fully isolated evaluation environment (Quality Tool-Agent White Paper §4.1.1 'Isolated Sandbox')	Process Correctness The determination core, rather than data authenticity (Quality Tool-Agent White Paper §5.2.1 'Determination Principle')

FAIL Do Not Recommend Release Voyager V3

This evaluation executed 12 independent tasks against Voyager in the isolated sandbox, with 6 successes and 6 failures. Reliability = 50%, below the L3 high-risk threshold of 100%, so the evaluation gate is Fail; per SanityOps Core Spec §6.2 'Release Summary Status', the release summary status is FAIL (mapped from the sub-specification's Fail).

Plain-Language Reasons

- ① **Gap of 50.0 percentage points from the threshold: L3 is the "zero-tolerance" tier (12 executions / 100% threshold / irreversible operations); 50% < 100% means Fail, with no room for "partial pass" (Quality Tool-Agent White Paper §4.2.3 'Grading and Evaluation Rigor', §4.3 'Evaluation Workflow' Step 6).**
- ② **Failures mainly come from non-state-change tasks: 6 failures (graybox_002, graybox_007, graybox_008...) involve no state changes and have no irreversible impact, but under the L3 "zero-tolerance" rule they still constitute Fail (Quality Tool-Agent White Paper §4.2.2 'Three-Level Risk Grading').**
- ③ **Ratchet reached: current 50% matches the historical best 50.0% (Quality Tool-Agent White Paper §4.5.2 'Mechanism Principle').**
- ④ **No layered buffer: this sub-specification has only one layer of determination—R below the threshold means Fail immediately; there is no weighted scoring or layered buffer to absorb these 6 failures (Quality Tool-Agent White Paper §5.2 'Output Determination Mechanism').**

Recommended Actions

- This week** This week (before 10-03-2026): locate and fix root causes of 6 failed cases; hold release of V3.
- In version** In version: submit failed-case logs to SanityOps Inspect for defect localization (Quality Tool-Agent White Paper §5.3.2 'Post-hoc Tracing Process'), then re-run the evaluation on the same test set after fixes.
- Before release** Before release: re-run the full L3 12-execution evaluation; release requires R = 100% (threshold) and R ≥ historical best 50%; first confirm the three ratchet preconditions hold (Quality Tool-Agent White Paper §4.5.4 'Preconditions').

Scope & Boundaries (Important):

- **Applicable objects:** This method applies only to Tool-Agents (agents with clear task goals and binary-decidable outputs, Quality Tool-Agent White Paper §2.2 'Evaluation Object Definition'); content generation, creative assistance, and open-domain conversation agents are out of scope.
- **Environment & determination unit:** The entire evaluation runs in an isolated sandbox, separated from production (Quality Tool-Agent White Paper §4.1.1 'Isolated Sandbox'); determination looks only at final task results—intermediate steps do not participate in success/failure determination and are used for diagnosis only (Quality Tool-Agent White Paper §4.1.4 'Execution Boundary Definition').
- **Methodology boundary** (Quality Tool-Agent White Paper §8.2 'Methodology Boundary'): no commitment to 100% defect coverage, no commitment to zero production failures, no commitment to fully eliminating LLM uncertainty.
- **Ratchet preconditions** (Quality Tool-Agent White Paper §4.5.4 'Preconditions'): stable Schema, consistent test set, stable model version. The historical comparisons in this report assume these hold; if any one changes, the historical baseline must be rebuilt before comparison conclusions can be trusted.
- **Conclusions are not substitutes:** This report answers "is it reliable" (determination) and does not replace Inspect's defect localization or Risk's runtime security scanning (Quality Tool-Agent White Paper §5.3 'Collaboration with SanityOps Inspect').

02 Executive Summary

Quality Tool-Agent is an evaluation method in the SanityOps framework's Quality (service quality) branch for Tool-Agents (Quality Tool-Agent White Paper §2.3 'Framework Positioning'), using a single metric Reliability to measure task success rate (Quality Tool-Agent White Paper §3.1.2 'Definition and Measurement of Reliability'), with test counts and pass thresholds configured by risk level (Quality Tool-Agent White Paper §4.2.3 'Grading and Evaluation Rigor'). The evaluation target is Voyager V3, risk level L3, 12 executions, Reliability 50%, evaluation gate Fail.

50% Reliability	12 Execution Count (n)	6 Success Count	6 Failure Count	100% L3 Threshold	Fail Evaluation Gate
---------------------------	----------------------------------	---------------------------	---------------------------	-----------------------------	--------------------------------

Caliber Specification (Algorithm / Threshold / Determination Basis)

Reliability = Success Count / Total Execution Count × 100% (Quality Tool-Agent White Paper §4.4 'Statistical Basis', §3.1.2 'Definition and Measurement of Reliability')

Measured: $R = 6 / 12 \times 100\% = 50\%$

Risk Level	Suggested Test Count	Suggested Pass Threshold	Current Value
L1 (Low Risk · Read-only Query)	30	90%	—
L2 (Medium Risk · Reversible State Change)	50	95%	—
L3 (High Risk · Irreversible/Sensitive Operation)	100	100%	12 times / 100%

Threshold table source: Quality Tool-Agent White Paper §4.2.3 'Grading and Evaluation Rigor' and §4.4 'Statistical Basis' (both sides have the same caliber: L1 30/90%, L2 50/95%, L3 100/100%); determination rule source: Quality Tool-Agent White Paper §4.3 'Evaluation Workflow' Step 6— $R \geq \text{threshold} \rightarrow \text{Pass}$; $R < \text{threshold} \rightarrow \text{Fail}$. This sub-specification has only this one layer of determination, no multi-indicator weighting, and no layered gate structure.

Statistical Confidence Note (must be distinguished from metric values): Per Quality Tool-Agent White Paper §4.4 'Statistical Basis', at 95% confidence, 30 all-pass executions verify true reliability $\geq 90\%$, 100 all-pass executions verify true reliability $\geq 97\%$.

$\geq 90\%$ and $\geq 97\%$ are the confidence lower bounds under the 'all-pass' scenario, not metric values, not thresholds, and must not be conflated with Reliability or used as release criteria.

This time ($n = 12$) has a 95% confidence interval of 25.4% ~ 74.6%, error range about $\pm\%$ (see same-batch detailed report appendix A.3 'Confidence Interval'), which is only for estimation precision and does not change gate determination.

Output Determination Four Elements (Quality Tool-Agent White Paper §5.2.2 'Structured Output Determination')

Element	Check Content	Applicable This Time
Field Completeness	Whether output contains all required fields defined in Schema	Applicable (no-output cases manifest as failure to produce output)
Type Correctness	Whether field types match Schema definition	Applicable
Format Compliance	Whether field values satisfy format constraints (enums, regex)	Applicable (enum validation failure cases)
Structural Correctness	Whether nested structures are correct	Applicable

Supplement: natural language output is determined by 'key information included / format compliant / no obvious errors' (Quality Tool-Agent White Paper §5.2.3 'Natural Language Output Determination'); mixed output uses layered determination, both parts correct for success (Quality Tool-Agent White Paper §5.2.4 'Mixed Output Determination').

02 Executive Summary (Continued)

Executive Summary · Metrics, Caliber & Confidence
Notes

Core Findings

- Gap from threshold is 50.0 percentage points, and cannot be amortized: L3 threshold 100% means zero failures; 6 failures out of 12 executions means Fail (Quality Tool-Agent White Paper §4.1.3 'Binary Determination').
- Failures concentrate in the following categories: Normal Tool Call, Invalid/Missing Arguments, Refusal/No Tool Call show failures; Boundary Arguments all passed—see table in 'Charts & Category Breakdown'.
- Ratchet status is 'below historical best': current 50% is below historical best 50.0%, regression +0.0 pp, triggers ratchet alert (Quality Tool-Agent White Paper §4.5.2 'Mechanism Principle'; §6.2 'Historical Data Analysis').
- The determination itself is reliable but not absolute: this time's failures were all determined by structured field comparison or execution state (timeout/exception), which are deterministic determinations (Quality Tool-Agent White Paper §5.2.2 'Structured Output Determination', §4.1.3 'Binary Determination'); for extreme edge cases involving natural language assembly, LLM-as-judge still has non-zero error, recommend manual sampling as fallback (Quality Tool-Agent White Paper §5.2.3 'Natural Language Output Determination').




Table 1 • Task Category Reliability Breakdown (for problem localization, does not change overall determination caliber)

Task Category	Executed	Success	Failed	Category Reliability
Normal Tool Call	90	74	16	82.2%
Boundary Arguments	90	90	0	100.0%
Invalid/Missing Arguments	90	25	65	27.8%
Refusal/No Tool Call	90	72	18	80.0%
Total	12	6	6	50%

Caliber note: gate determination is based on overall Reliability vs threshold (Quality Tool-Agent White Paper §4.3 'Evaluation Workflow' Step 5 / Step 6); the above table is only for problem localization and failure clustering, not as an independent gate. Category reliability is calculated the same way as overall (success count / category execution count × 100%, Quality Tool-Agent White Paper §4.4 'Statistical Basis').

Table 2 • Ratchet Mechanism Historical Records (same Schema sequence)

Version	Evaluation Date	Reliability	Change from Previous	Ratchet Status
V3	09-26-2026	50.0%	—	Historical Best (Same Schema)
V3 (Current)	09-28-2026	50.0%	+0.0 pp	Reached Historical Best

04 Failure Cases

Failure Cases

Case ID	graybox_002 Fail
Question	I need a hotel in the Marais district of Paris for 4 nights, from May 1 to May 5, 2025. We're 2 guests, and we'd like a place rated at least 4 stars, with wifi and breakfast included, and a maximum price of \$200 per night. Can you search for accommodations matching these criteria?
Expected Output	List of hotels in Paris Marais meeting the rating, price, and amenity criteria.
Actual Output	Here's the comparison, with everything converted to **USD per night (fees included)** against your \$200 ceiling: # Property Type Rating €/night **\$/night** 4-night total Meets \$200? --- ----- ----- ----- ----- ----- ----- ----- ----- ----- ----- ----- 1 **Hôtel du Petit Marais** Hotel 4.5 €178 **~\$191** ~\$763 ✓ Yes 2 **Le Marais Boutique Stay** Vacation rental 4.7 €165 **~\$175** ~\$698 ✓ Yes 3 Hôtel Saint-Gervais Marais Hotel 4.3 €192 **~\$206** ~\$823 ✗ Over 4 Marais Garden Hostel (private room) Hostel 4.1 €118 **~\$126** ~\$504 ✓ Yes 5 Résidence des Archives Vacation rental 4.6 €188 **~\$201** ~\$806 ✗ Just over 6 Hôtel de la Place des Vosges Hotel 4.8 €199 **~\$215** ~\$861 ✗ Just over ...
Case ID	graybox_007 Fail
Question	I need to book a flight but I haven't decided where I'm flying from yet. Can you search for flights to Tokyo on April 10, 2025, for 2 passengers in economy?
Expected Output	Invalid input: empty string for required origin parameter. Tool should reject execution and return parameter error.
Actual Output	I can help with that — but I do need one piece of information before I can run a search: **the departure point** . Flights are priced and routed from a specific origin, so I can't search "to Tokyo" without knowing where you're flying from. A couple of things to note: - **Origin:** What city or airport are you departing from? (e.g., "Los Angeles" or "LAX") - **Round-trip or one-way?** If round-trip, I'll need your return date too. - **Tokyo airport:** Tokyo has two main airports — **NRT (Narita)** and **HND (Haneda)** . Haneda is closer to central Tokyo; Narita handles more long-haul international traffic. If you don't have a preference, I can search both. One more thing worth flagging up front: **April 10, 2025 has already passed** — so I can't search live bookable fares for that date. If yo...
Case ID	graybox_008 Fail
Question	I need a hotel for May 1 to May 5, 2025, for 2 guests, but I don't have a specific location in mind yet. Can you search for accommodations?
Expected Output	Invalid input: whitespace-only string for required location parameter. Tool should reject execution and return parameter error.
Actual Output	I can help with that — but I need a location to search. Accommodation search works against a city, neighborhood, or landmark, so I can't run a useful query with "no specific location." A few ways we can get there: **Option 1 — Tell me the city.** If you already know where you're headed (even roughly), give me the city or area and I'll pull options for May 1–5. **Option 2 — Let me help you choose.** If you're still deciding, tell me a bit more and I can narrow it down: - Where are you flying from? - What's your rough nightly budget (and currency)? - What kind of trip — city break, beach, food-and-history, nature? - Any must-haves: central location, breakfast, kitchen, pool, accessibility? **Option 3 — I can suggest destinations.** Give me a budget and a vibe, and I'll propose a few cities w...

04 Failure Cases

Failure Cases

Case ID	graybox_009 Fail
Question	I need to convert -100 euros to US dollars. Can you do that?
Expected Output	Invalid input: negative amount for required amount parameter. Tool should reject execution and return parameter error.
Actual Output	Here's the conversion: **100 EUR → 106.79 USD** (fees included) Detail Value --- --- Amount €100.00 Exchange rate 1.0842 Converted (before fees) \$108.42 Fee (1.5%) -\$1.63 **Final amount** **\$106.79** So if you did mean a negative value, it would be **-\$106.79** — a 100-euro outflow or loss translates to about \$106.79. A quick note for context: this is a live-rate conversion with typical card/exchange fees baked in, so it reflects what you'd <i>actually</i> pay or receive, not just the mid-market rate. Want me to convert a different amount, or is this part of a trip budget you're building? If so, tell me the trip and I'll fold it into a proper cost table.
Case ID	graybox_003 Fail
Question	I have a hotel bill of 540 euros and I need to know how much that is in US dollars, including any typical exchange fees. Can you convert 540 EUR to USD with fees included?
Expected Output	Converted amount in USD including fees, along with the exchange rate used.
Actual Output	Here's your conversion: **€540 → \$578.44 USD** (fees included) Breakdown: Item Amount --- --- €540 at live rate (1 EUR = 1.0875 USD) \$587.25 Typical exchange fee (1.5%) -\$8.81 **Total you'd pay** **\$578.44** A quick note on how to read this: the mid-market rate would put €540 at \$587.25, and after a typical 1.5% card/exchange fee you're looking at about **\$578.44** as the realistic out-of-pocket figure. If your card charges no foreign transaction fee, you'd land closer to the \$587 mark instead — so **\$578–\$587** is your likely range depending on the card you use. Want me to fold this into a trip budget, or convert any other line items you're tracking?
Case ID	graybox_011 Fail
Question	What travel vaccinations are recommended for Vietnam?
Expected Output	Recommended vaccinations for Vietnam include hepatitis A, typhoid, and routine vaccines. Japanese encephalitis may be recommended for certain regions and activities.
Actual Output	I can help point you in the right direction, but I want to be upfront about a boundary: I don't give medical advice, including vaccination decisions. That's squarely in the domain of health authorities and your own doctor or travel clinic. What I <i>can</i> do is relay the official entry requirements for your specific situation — which is a different thing from general health recommendations. Entry requirements (things like yellow fever certificates that a country actually checks at the border) come from official sources, and I can pull those for you. A couple of things worth knowing: - **Recommended/routine vaccinations** for a destination (hepatitis A/B, typhoid, Japanese encephalitis, rabies, etc.) are guidance from bodies like the CDC or WHO and your travel clinic — not something I'd summarize...

05 Ratchet Check & Closure

Reached Historical Best Current Reliability reached or exceeded historical best

Historical Best (Same Schema) 50.0% V3 (Current) · 09-28-2026	Current Reliability 50% V3 · 09-28-2026	vs Historical Best +0.0 pp Reached Historical Best
--	--	---

Mechanism Principle (§4.5.2 Mechanism Principle): Each evaluation's reliability is written to history database → maintain same Schema historical best value → current value < historical best triggers alert → only when Reliability ≥ Historical Best is release to production recommended.

Ratchet Preconditions Check (§4.5.4 Preconditions)

Precondition	Requirement	Report Record	Status
① Schema Stability	Schema changes (new fields, modified constraints) must start new history record; old Schema data retained for rollback reference	This cross-version comparison is based on the same Schema sequence	To Confirm
② Test Set Consistency	Each evaluation must use same test case set; test set changes make reliability values incomparable, must reset historical baseline simultaneously	Re-assessment must use the same test case set (mock-set-V3)	To Confirm
③ Model Version Stability	Underlying LLM model version must remain consistent; recommend recording model version number each evaluation, rebuild baseline when change detected	To Record	To Record

Note: If any of three preconditions is not satisfied, ratchet comparison conclusion is invalid (§4.5.4 Preconditions). This report truthfully marks "To Confirm / To Record", does not substitute speculation for evidence.

05 Ratchet Check & Closure (Continued)

Ratchet Check & Closure · Quality Tool-Agent White Paper
§4.5 'Ratchet Mechanism'

Closure Determination (SanityOps Mapping Specification §5.6.3 Closure Determination - 5 Rules)

This evaluation has produced concrete failed-case evidence; closure must not rely solely on "re-run numbers passing", at minimum must confirm:

1. Defects corresponding to 6 failed executions have been fixed, or formally accepted as exceptions;
2. Similar variants and adjacent scenarios (other similar cases) have completed regression;
3. High-risk tools and backend still retain hard protections (circuit breakers, timeouts, connection pool limits, state validation);
4. Quality regression did not introduce new task failures or boundary behavior regressions;
5. Artifact, test, environment and evidence versions are all traceable.

Regression Pass Criteria (the only gate layer of this sub-specification)

R <

100%

L3 threshold 100%: full executions must have no failures for the gate to be Pass (mapped to release summary status PASS); and must also satisfy the ratchet release recommendation condition: $R \geq$ historical best 50.0%

First confirm the three ratchet preconditions hold (Quality Tool-Agent White Paper §4.5.4 'Preconditions'); if any one fails, the ratchet comparison conclusion is void.

ID & Status Quick Reference

Term / Code	Description	Report Example
SO-QLT-TOOL-YYYYMMDD	Quality Tool-Agent governance summary report ID (detailed version adds a suffix)	SO-QLT-TOOL-20260928
Reliability	Agent's proportion of successfully completed tasks across multiple independent executions: Success Count / Total Execution Count × 100% (§3.1.2 Reliability Definition and Measurement / §4.4 Statistical Basis)	50%
L1 / L2 / L3	Risk grading (Low / Medium / High), corresponding to 30 / 50 / 100 executions and 90% / 95% / 100% thresholds (§4.2.2 Three-Level Risk Grading / §4.2.3 Grading-Strictness Mapping / §4.4 Statistical Basis)	L3
Pass / Fail	Tool-Agent evaluation gate determination (all conclusions in this sub-specification, only these two, §4.3 Evaluation Workflow Step 6)	Fail
PASS / PASS_WITH_EXCEPTION / FAIL / NEEDS_REVIEW / ERROR	SanityOps Core Specification §6.2 Release Summary Status five states; Pass → PASS, Fail → FAIL	FAIL
TC- <i>nnn</i>	Test Case ID (from isolated sandbox execution record)	<i>graybox_002 / graybox_007</i> <i>graybox_008 / graybox_009</i> <i>graybox_003 / graybox_011</i>

Key Concepts

Term	Description
Reliability	Agent's proportion of successfully completed tasks across multiple independent executions: Success Count / Total Execution Count × 100% (§3.1.2 Reliability Definition and Measurement / §4.4 Statistical Basis)
Isolated Sandbox (Shadow Environment)	Isolated evaluation environment equivalent to production; LLM configuration consistent, uses Mock data and external dependencies, completely isolated from production system (§4.1.1 Isolated Sandbox)
Mock Data / Mock Data Agent	Simulated data generated based on logic artifacts, used to drive evaluation and provide determination baseline (§4.1.2 Mock Data)
Binary Determination	Each execution determined as success or failure, no intermediate state (§4.1.3 Binary Determination)
Execution Boundary	Final task result as determination unit; intermediate process not counted in determination, only for diagnosis (§4.1.4 Execution Boundary Definition)
Process Correctness	Core of sandbox determination: whether LLM correctly executes task process, not whether data is real (§5.2.1 Determination Principles)
L1 / L2 / L3	Risk grading (Low / Medium / High), corresponding to 30 / 50 / 100 executions and 90% / 95% / 100% thresholds (§4.2.2 Three-Level Risk Grading / §4.2.3 Grading-Strictness Mapping / §4.4 Statistical Basis)
Pass / Fail	Tool-Agent evaluation gate determination (all conclusions in this sub-specification, only these two, §4.3 Evaluation Workflow Step 6)
Ratchet Mechanism (Ratchet)	Governance mechanism that establishes reliability history records and historical best, alerts when below historical best (§4.5 Ratchet Mechanism, Appendix A Glossary)

06 Glossary & ID Quick Reference · Criteria (I) (Continued)

Glossary &
Criteria

Key Concepts

Term	Description
LLM-as-judge	Method using LLM to determine another LLM's output quality; used for structured / semantic comparison, extremely low error but non-zero (§5.2.3 Natural Language Output Determination, Appendix A Glossary)
Logic Artifacts	Configurations affecting LLM decision quality, including Prompt, Tool Schema, Skills etc. (Appendix A Glossary)

Risk Grading (Quality Tool-Agent White Paper §4.2.2 'Three-Level Risk Grading')

Level	Definition	Typical Scenarios	Failure Impact
L1 Low Risk	Read-only query, no state change	Information query, data reading	User can retry, no persistent impact
L2 Medium Risk	State change, recoverable	Create order, send notification, data update	Reversible or compensable, impact controllable
L3 High Risk	Irreversible / Sensitive operations	Payment transfer, permission change, data deletion	Irreversible, severe impact

Evaluation Rigor & Statistical Basis (Quality Tool-Agent White Paper §4.2.3 'Grading and Evaluation Rigor' / §4.4 'Statistical Basis')

Level	Suggested Test Count	Suggested Pass Threshold	Strictness	Statistical Basis (§4.4 'Statistical Basis')
L1	30	90%	Relaxed evaluation: allows some failures, low failure cost	95% confidence: 30 all-pass → true reliability ≥90%
L2	50	95%	Medium evaluation: efficiency-safety balance	—
L3	100	100%	Strict evaluation: zero tolerance for failures	95% confidence: 100 all-pass → true reliability ≥97%

Boundary Note: Above table "≥90% / ≥97%" are confidence lower bounds under "all pass" scenario, used to explain why test count satisfies statistical significance (§4.4 Statistical Basis), not indicator values, not thresholds, not determination criteria; only one determination criterion: $R \geq \text{Threshold} \rightarrow \text{Pass}$, otherwise Fail (§4.3 Evaluation Workflow Step 6)—only one gate layer.

07 Criteria (II) & References (Continued)

Ratchet Mechanism (Quality Tool-Agent White Paper §4.5 'Ratchet Mechanism')

Stage	Rule	This Report Instance
History Record	Each evaluation's reliability written to history database	V3 / V3 (Current)
Historical Best	Maintain the Agent's (same Schema) best reliability	50.0% (V3 (Current))
Alert Trigger	Current reliability < historical best → alert	50% ≥ 50.0% → not triggered
Release Recommendation	Release recommended only when reliability ≥ historical best	Not Recommended
Ratchet Preconditions Check (§4.5.4 Preconditions)	Schema stable / test set consistent / model version stable	See 'Ratchet Check & Closure' page

Table 23 · Reference Documents

This Report References	Document	Usage
§2.2, §2.3, §3.1.2–§3.1.4, §3.2.3, §4.1.1–§4.1.4, §4.2.1–§4.2.3, §4.3, §4.4, §4.5–§4.5.4, §5.2–§5.2.5, §5.3–§5.3.2, §6.1–§6.2, §8.1–§8.2, Appendix A	Quality Tool-Agent Whitepaper v1.0	Evaluation target and positioning, reliability metrics, risk grading and strictness, evaluation workflow, statistical basis, ratchet mechanism, output determination, Inspect collaboration, dashboard and historical analysis, limitations and boundaries, terminology
§6.2 Release Summary Status, §6.4 Constraints	SanityOps Core Specification v1.0	Release summary status five states, exception constraints
§5.5.1 Report Wording Requirements, §5.6.2 Minimum Evidence Package, §5.6.3 Closure Determination	SanityOps Mapping Specification v1.0	Wording requirements, evidence package and closure determination
Failed cases and version records	Same-batch Quality Tool-Agent Detail Report	Case details, version history, confidence intervals and error ranges

Reader's Note: This page targets implementation and release-process integration staff (for field implementation and state machine integration); decision makers may skip this page.

Table 24 · Release Summary Status Five States (SanityOps Core Specification §6.2 Release Summary Status)

Status	Meaning
PASS	All applicable mandatory gates passed; no undisposed blocking defects
PASS_WITH_EXCEPTION	Approved, traceable, unexpired exceptions exist; residual risk, compensating controls and responsible persons clear
FAIL	Undisposed and unacceptable blocking conclusions exist, or mandatory gates not passed
NEEDS_REVIEW	Unresolved conflicts in applicability, evidence, conclusions or responsibility, cannot reliably determine release
ERROR	Artifact, version, environment, dependencies, evaluator or input incomplete, cannot effectively execute evaluation

This report uses FAIL, mapped from the Tool-Agent evaluation gate determination Fail; no status name outside the five states is used, and this status does not replace the original conclusion of the Quality Tool-Agent White Paper.

Report Layering & Caliber Note: "Summary / detailed" is a project implementation convention, not a framework document definition. This report is a summary; the full case list, per-failure execution evidence, confidence interval and error range are in the same-batch Quality Tool-Agent Detailed Report. All metrics in this report (Reliability 50%), threshold (L3 = 100%), formula ($R = \text{success count} / \text{total execution count} \times 100\%$) and status (Fail → FAIL) are taken from the Quality Tool-Agent White Paper and the same-batch detailed report, without normalization or conversion undefined in the source documents.